

# Emergency Room Management Problems

---

# Emergency Room @ Maggiore Hospital

**We will now consider a couple of Emergency Room management problems**

We will use a dataset for the "Maggiore" hospital in Bologna



We will consider two main use cases:

- Predicting future arrivals (a Machine Learning problem)
- Managing the ER center resources (a Combinatorial Optimization problem)

# A Look at the Dataset

We will start as usual by having a look at the dataset

```
In [2]: data = er.load_data(data_folder)
data
```

Out [2]:

	year	ID	Triage	TkCharge	Code	Outcome	Flow
0	2018	1	2018-01-01 00:17:33	2018-01-01 04:15:36	green	admitted	[triage,visit,RX,visit]
1	2018	2	2018-01-01 00:20:33	2018-01-01 03:14:19	green	admitted	[triage,visit,lab,visit]
2	2018	3	2018-01-01 00:47:59	2018-01-01 04:32:30	white	admitted	[triage,visit,otolaryngological visit,visit]
51238	2018	51239	2018-01-01 00:49:51	NaT	white	abandoned	[]
51240	2018	51241	2018-01-01 01:00:40	NaT	green	abandoned	[]
...	...	...	...	...	...	...	...
95665	2019	95666	2019-10-31 23:26:54	2019-10-31 23:41:13	yellow	admitted	[triage,visit,RX,visit]
95666	2019	95667	2019-10-31 23:46:43	2019-11-01 09:30:25	green	admitted	[triage,visit]
108622	2019	108623	2019-10-31 23:54:05	NaT	green	abandoned	[]
95667	2019	95668	2019-10-31 23:55:32	2019-11-01 00:18:46	yellow	admitted	[triage,visit]
108623	2019	108624	2019-10-31 23:59:21	NaT	green	abandoned	[]

108625 rows × 7 columns

# A Look at the Dataset

```
In [4]: data.iloc[:3]
```

Out[4]:

	year	ID	Triage	TkCharge	Code	Outcome	Flow
0	2018	1	2018-01-01 00:17:33	2018-01-01 04:15:36	green	admitted	[triage,visit,RX,visit]
1	2018	2	2018-01-01 00:20:33	2018-01-01 03:14:19	green	admitted	[triage,visit,lab,visit]
2	2018	3	2018-01-01 00:47:59	2018-01-01 04:32:30	white	admitted	[triage,visit,otolaryngological visit,visit]

- Each row refers to a single patient
- `Triage` is the arrival time of each patient
- `TkCharge` is the time when a patient starts the first visit
- `code` refers to the estimated priority (white < green < yellow < red)
- `outcome` discriminates some special conditions (people quitting, fast tracks)
- `Flow` is the sequence of treatments that actually took place

# A Look at the Dataset

Let's also have a look at the data types

```
In [5]: data.dtypes
```

```
Out[5]: year                int64  
        ID                  int64  
        Triage              datetime64[ns]  
        TkCharge             datetime64[ns]  
        Code                 category  
        Outcome              category  
        Flow                 object  
        dtype: object
```

■ Flow is actually a string

**We will initially focus for now on predicting arrivals**


...Hence, it makes sense to sort rows by increasing triage time:

```
In [8]: data.sort_values(by='Triage', inplace=True)
```

# Inter-Arrival Times

Let's check empirically the distribution of the inter-arrival times

```
In [6]: f = '2min'
delta = data['Triage'].iloc[1:].round(f) - data['Triage'].iloc[:-1].round(f).values
er.plot_bars(delta.value_counts().sort_index(), tick_gap=10, figsize=figsize)
```




- There is a number of very low inter-arrival times
- It may be due to how triage is performed (bursts, rather than a steady flow)

# Waiting Time

Here is the distribution of the waiting times

```
In [7]: tmp = data[~data['TkCharge'].isnull()]\nwait_time = tmp['TkCharge'].round('10min') - tmp['Triage'].round('10min')\ner.plot_bars(wait_time.value_counts().sort_index(), tick_gap=10, figsize=figsize)
```




- The distribution is heavy-tailed
- I.e. the probability of very long waiting times is non-negligible

# Code Distribution

## The distribution of the priority codes

```
In [8]: er.plot_bars(data['Code'].value_counts(), figsize=figsize)
```




- Green code (low severity) form the majority of arrivals
- Yellow and red codes (mid and high severity) are in smaller numbers
- White codes (lowest priority) are also not very frequent



# Outcome Distribution

```
In [9]: er.plot_bars(data['Outcome'].value_counts(), figsize=figsize)
```




- Abandons are infrequent, as are "fast track" patients

# Arrival Distribution over Months

Let's look at the arrival distribution over months

```
In [10]: months = data['Triage'].dt.month  
er.plot_bars(months.value_counts().sort_index(), figsize=figsize)
```




- The low values for Nov. and Dec. are due to the 2019 series ending in October
- The distribution seems stable (but we are not plotting standard deviations!)

# Arrival Distribution over Weekdays

Let's look at the distribution over weekdays

```
In [11]: weekdays = data['Triage'].dt.weekday  
er.plot_bars(weekdays.value_counts().sort_index(), figsize=figsize)
```




- Similarly to months, weekdays are likely to have little predictive power
- ...But it's better not to rush conclusions (we still are not plotting the stddev!)

# Arrival Distribution over Hours

Let's see now the arrival distribution over the hours of the day

```
In [12]: hours = data['Triage'].dt.hour  
er.plot_bar(hours.value_counts().sort_index(), figsize=figsize)
```



- There is a clear pattern: the hour of the day will have strong predictive power
- Again, analyzing the standard deviation may provide better insights

# Data Preparation

---

# Binning

## In our two considered problems:

- We are not going to revise our decisions continuously
- We are not interested in predicting the next arrival

Rather:

- We will take decisions **at fixed intervals**
- We care about the expected arrivals **in a given horizon**

**Overall, we need to choose a meaningful time unit**

In other words, we need to perform some kind of binning

- We used binning to downsample high-frequency data
- ...We will use binning to **aggregate events with a variable frequency**

# Code-Based Counts

We will prepare the data to track counts for all priority codes

```
In [13]: codes = pd.get_dummies(data['Code'])
codes.set_index(data['Triage'], inplace=True)
codes.columns = codes.columns.to_list()
print(f'Number of examples: {len(codes)}')
codes.head()
```

Number of examples: 108625

Out[13]:

	green	red	white	yellow
Triage				
2018-01-01 00:17:33	1	0	0	0
2018-01-01 00:20:33	1	0	0	0
2018-01-01 00:47:59	0	0	1	0
2018-01-01 00:49:51	0	0	1	0
2018-01-01 01:00:40	1	0	0	0

- The `get_dummies` function applies a one-hot encoding to categorical value
- The method generates a categorial column index (then converted to list)

# Resampling

Then, we need to aggregate data with a specified frequency

```
In [16]: codes_b = codes.resample('H').sum()
print(f'Number of examples: {len(codes_b)}')
codes_b.head()
```

Number of examples: 16056

Out[16]:

	green	red	white	yellow
Triage				
2018-01-01 00:00:00	2	0	2	0
2018-01-01 01:00:00	7	1	1	1
2018-01-01 02:00:00	4	1	4	3
2018-01-01 03:00:00	7	0	1	1
2018-01-01 04:00:00	3	0	2	0

- We used the `resample` iterator
- `resample` generates a dataframe with a **dense** index
- We chose 1 hour as our time unit



# Computing Totals

We also compute the total number of arrivals for each interval

```
In [18]: cols = ['white', 'green', 'yellow', 'red']  
codes_b['total'] = codes_b[cols].sum(axis=1)  
codes_b.head()
```

Out[18]:

	green	red	white	yellow	total
Triage					
2018-01-01 00:00:00	2	0	2	0	4
2018-01-01 01:00:00	7	1	1	1	10
2018-01-01 02:00:00	4	1	4	3	12
2018-01-01 03:00:00	7	0	1	1	9
2018-01-01 04:00:00	3	0	2	0	5


- We use a fixed column list to obtain the same result for multiple executions
- The total count will be less noisy, if the individual terms are independent

# Counts over Time

**Our resampled series can be plotted easily over time**

Let's see the total counts as an example:

```
In [19]: er.plot_series(codes_b['total'], figsize=figsize)
```




# Counts over Time

**Our resampled series can be plotted easily over time**

The same plot, for the red codes (the counts are significantly lower):

```
In [20]: er.plot_series(codes_b['red'], figsize=figsize)
```




**We are now ready to tackle the first of our ER management problems**

# Variability

**With our binned series, we can compute standard deviations**

Let's check it over different months:

```
In [21]: means = codes_b.groupby(codes_b.index.month).mean()  
stds = codes_b.groupby(codes_b.index.month).std()  
er.plot_series(means['total'], std=stds['total'], figsize=figsize)
```




- The variability does not change much over different months

# Variability

Here is the standard deviation over weekdays

```
In [22]: means = codes_b.groupby(codes_b.index.weekday).mean()  
stds = codes_b.groupby(codes_b.index.weekday).std()  
er.plot_series(means['total'], std=stds['total'], figsize=figsize)
```




■ A decreasing trend, but rather weak

# Variability

## ...And finally over hours

```
In [25]: means = codes_b.groupby(codes_b.index.hour).mean()  
stds = codes_b.groupby(codes_b.index.hour).std()  
er.plot_series(means['total'], std=stds['total'], figsize=figsize)
```



- Again, unlike the mean, the stdev is similar over the hours of a day