

Arrival Prediction

Arrival Prediction

We can now frame our arrival prediction problem

We want to predict the number of arrivals in the next interval

- We will focus on predicting the **total number of arrivals**
- The same models can be applied to any of the individual counts

Technically, this is a regression problem

- ...Which does **not** mean that an MSE is the best choice
- It makes more sense to **check the target distribution** first

Main issue: our regressor will learn a conditional probability distribution

- ...So, that's the kind of distribution that we should check, in principle
- It is a difficult task, since we do not know yet which input we are going to use

However, we already know that **some features** are likely to be useful


Analyzing the Conditional Arrival Distribution

...In particular, we know that the hour of the day is a good predictor

Let's check the (conditional) distribution for 6am:

```
In [37]: tmp = codes_b[codes_b.index.hour == 6]['total']  
tmpv = tmp.value_counts(sort=False, normalize=True).sort_index()  
er.plot_bars(tmpv, figsize=figsize)
```

Figure



■ This is **not** a normal distribution

Poisson Distribution

When we need to count occurrences over time...

It's almost always worth checking the Poisson distribution, which models:

- The number of occurrences of a certain event in a given interval
- ...Assuming that these events are independent
- ...And they occur at a constant rate

In our case:

- The independence assumption is reasonable (arrivals do not affect each other)
- The constant rate is true for the conditional probability
- ...Assuming that our predictor is good enough

Intuitively: the predictor will estimate the constant rate

Poisson Distribution

The Poisson distribution is defined by a single parameter λ

λ is the rate of occurrence of the events

- The distribution has a **discrete support**
- The Probability Mass Function is:

$$p(k, \lambda) = \frac{\lambda^k e^{-\lambda}}{k!}$$


- Both the **mean** and the **standard deviation** have the same value (i.e. λ)
- The distribution skewness is $\lambda^{-\frac{1}{2}}$
 - For low λ values, there is a significant positive skew (to the left)
 - The distribution becomes less skewed for large λ

Fitted Poisson Distribution

Let's try to fit a Poisson distribution over our target

```
In [39]: mu = tmp.mean()
dist = stats.poisson(mu)
x = np.arange(tmp.min(), tmp.max()+1)
er.plot_bar(tmpv, figsize=figsize, series=pd.Series(index=x, data=dist.pmf(x)))
```

Figure




It's a very good match!

Fitted Poisson Distribution

Let's try for 8AM (closer to the peak)

```
In [41]: tmp = codes_b[codes_b.index.hour == 8]['total']  
tmpv = tmp.value_counts(sort=False, normalize=True).sort_index()  
mu = tmp.mean()  
dist = stats.poisson(mu)  
x = np.arange(tmp.min(), tmp.max()+1)  
er.plot_bars(tmpv, figsize=figsize, series=pd.Series(index=x, data=dist.pmf(x)))
```

Figure



Fitted Poisson Distribution

...And finally for the peak itself

```
In [47]: tmp = codes_b[codes_b.index.hour == 11]['total']
tmpv = tmp.value_counts(sort=False, normalize=True).sort_index()
mu = tmp.mean()
dist = stats.poisson(mu)
x = np.arange(tmp.min(), tmp.max()+1)
er.plot_bars(tmpv, figsize=figsize, series=pd.Series(index=x, data=dist.pmf(x)))
```

Figure




Table Model

Rate Table

We know that:

- A Poisson distribution is a good fit for our target...
- ...But the rate depends (at least) on the hour of the day

The simplest way to take this into account is using a **lookup table**

- The table will contain average arrival values for each hour of the day
- ...But first, we need to separate the training and test data

```
In [49]: sep = '2019-01-01'  
tr_data = codes_b[codes_b.index < sep]  
ts_data = codes_b[codes_b.index >= sep]
```

Rate Table

We can then build our "rate table"

```
In [51]: rate_table = tr_data.groupby(tr_data.index.hour).mean()  
rate_table.head()
```

Out [51]:

	green	red	white	yellow	total
Triage					
0	2.265753	0.219178	0.367123	1.054795	3.906849
1	1.915068	0.213699	0.331507	0.860274	3.320548
2	1.734247	0.202740	0.287671	0.780822	3.005479
3	1.515068	0.153425	0.232877	0.687671	2.589041
4	1.334247	0.134247	0.197260	0.717808	2.383562

- We are computing (average) rates for all the codes
- This will enable changing focus to a different target, if we so wish

Predictions

We can not obtain the predictions

- We need to associate each example with the correct rates
- The rates themselves are then the predictions

```
In [52]: def preds_from_rate_table(data, rate_table):  
    tmp = data.copy()  
    tmp['hour'] = data.index.hour  
    tmp = tmp.join(rate_table, on='hour', lsuffix='_orig')  
    return tmp[rate_table.columns]  
  
tr_preds = preds_from_rate_table(tr_data, rate_table)  
ts_preds = preds_from_rate_table(ts_data, rate_table)
```

- We use a `join` operation to associate examples and rates

Evaluation


We can now evaluate the predictions. First for the **training set**:

```
In [53]: slabel = 'total'
er.plot_pred_scatter(tr_preds[slabel], tr_data[slabel], figsize=figsize)
print(f'R2: {metrics.r2_score(tr_data[slabel], tr_preds[slabel])}')
print(f'MAE: {metrics.mean_absolute_error(tr_data[slabel], tr_preds[slabel])}')
```

R2: 0.5940116186135868

MAE: 1.945866016138112

Figure




Evaluation

...And then for the test set:

```
In [54]: er.plot_pred_scatter(ts_preds['total'], ts_data['total'], figsize=figsize)
print(f'R2: {metrics.r2_score(ts_data[slabel], ts_preds[slabel])}')
print(f'MAE: {metrics.mean_absolute_error(ts_data[slabel], ts_preds[slabel])}')
```

R2: 0.5878063024665628
MAE: 1.9621410117760154



These results will be our baseline

Probabilistic Neural Model

Probabilistic Neural Model

We will now try to learn a hybrid probabilistic-neural model

Using a neural network makes adding inputs easier

- We know that there is a slight decreasing trend along week days
- ...So let's try adding the week day as a numerical input

```
In [65]: tr_data_in = keras.utils.to_categorical(tr_data.index.hour)
tr_data_in = np.hstack((tr_data_in, (tr_data.index.weekday/7).values.reshape(-1,1) ))

ts_data_in = keras.utils.to_categorical(ts_data.index.hour.values)
ts_data_in = np.hstack((ts_data_in, (ts_data.index.weekday/7).values.reshape(-1,1)))
```

- We use a categorical encoding for the day hour, due to its non-linear effect
- ...But it is (mostly) fine to use the week day as number, due to its linear effect

The Architecture

First we define our architecture

```
In [66]: def build_probabilistic_regressor(input_shape, hidden):
          model_in = keras.Input(shape=input_shape, dtype='float32')
          x = model_in
          for h in hidden:
              x = layers.Dense(h, activation='relu')(x)
          log_rate = layers.Dense(1, activation='linear')(x)
          lf = lambda t: tfp.distributions.Poisson(log_rate=t, force_probs_to_zero_outside_support=False)
          model_out = tfp.layers.DistributionLambda(lf)(log_rate)
          model = keras.Model(model_in, model_out)
          return model

          nn2 = build_probabilistic_regressor(tr_data_in.shape[1], hidden=[32])
```

- We use an MLP, plus a `DistributionLambda` layer for the Poisson distribution
- We will try with a shallow network

The Architecture

We can also plot the architecture

```
In [67]: nn2.summary()
```

```
Model: "model_3"
```

Layer (type)	Output Shape	Param #
=====		
input_4 (InputLayer)	[(None, 25)]	0
dense_6 (Dense)	(None, 32)	832
dense_7 (Dense)	(None, 1)	33
distribution_lambda_3 (Distr multiple		0
=====		

```
Total params: 865  
Trainable params: 865  
Non-trainable params: 0
```

- More parameters than the table (and initially it will work worse)

Training

We train the model for maximum likelihood, as usual

```
In [68]: negloglikelihood = lambda y_true, dist: -dist.log_prob(y_true)
nn2.compile(optimizer='Adam', loss=negloglikelihood)
cb = [callbacks.EarlyStopping(patience=10, restore_best_weights=True)]
history2 = nn2.fit(tr_data_in, tr_data['total'].values.astype(np.float32),
                  validation_split=0.2, callbacks=cb, batch_size=32, epochs=50, verbose=1)
```

```
Epoch 1/50
219/219 [=====] - 1s 1ms/step - loss: 4.4935 - val_loss: 2.4800
Epoch 2/50
219/219 [=====] - 0s 710us/step - loss: 2.3570 - val_loss: 2.2709
Epoch 3/50
219/219 [=====] - 0s 720us/step - loss: 2.2711 - val_loss: 2.2778
Epoch 4/50
219/219 [=====] - 0s 733us/step - loss: 2.2643 - val_loss: 2.2699
Epoch 5/50
219/219 [=====] - 0s 719us/step - loss: 2.2639 - val_loss: 2.2661
Epoch 6/50
219/219 [=====] - 0s 726us/step - loss: 2.2618 - val_loss: 2.2677
Epoch 7/50
219/219 [=====] - 0s 710us/step - loss: 2.2620 - val_loss: 2.2683
Epoch 8/50
219/219 [=====] - 0s 718us/step - loss: 2.2610 - val_loss: 2.2710
Epoch 9/50
219/219 [=====] - 0s 721us/step - loss: 2.2594 - val_loss: 2.2669
Epoch 10/50
219/219 [=====] - 0s 718us/step - loss: 2.2597 - val_loss: 2.2688
Epoch 11/50
```


Training

We check the loss behavior over time

```
In [69]: er.plot_training_history(history2, figsize=figsize)
tr2, vl2 = history2.history["loss"][-1], np.min(history2.history["val_loss"])
print(f'Loss: {tr2:.4f} (training, final), {vl2:.4f} (validation, best)')
```

Loss: 2.2576 (training, final), 2.2632 (validation, best)

Figure




Evaluation

We can now obtain and evaluate the predictions. On the **training set**:

```
In [70]: tr_preds2 = nn2(tr_data_in).mean()
slabel = 'total'
er.plot_pred_scatter(tr_preds2, tr_data[slabel], figsize=figsize)
print(f'R2: {metrics.r2_score(tr_data[slabel], tr_preds2)}')
print(f'MAE: {metrics.mean_absolute_error(tr_data[slabel], tr_preds2)}')
```

```
R2: 0.6059787926961983
MAE: 1.9191851117839551
```

Figure



Evaluation

And we get similar results for the test set

```
In [71]: ts_preds2 = nn2(ts_data_in).mean()
er.plot_pred_scatter(ts_preds2, ts_data['total'], figsize=figsize)
print(f'R2: {metrics.r2_score(ts_data[slabel], ts_preds2)}')
print(f'MAE: {metrics.mean_absolute_error(ts_data[slabel], ts_preds2)}')
```

R2: 0.6006004538247001
MAE: 1.9296061691681021

