

**Waiting Time**

---

# Waiting Time

So far, we have assumed the makespan as our cost metric

In an ER center, we are likely more interested in the waiting time

- This can be defined simply as the sum of the start times of the initial tasks

$$w = \sum_{i=1}^n s_i^{first}$$

- Where  $s_i^{first}$  is the start time of the first task of patient  $i$

**Patients are color-coded: it makes sense to track different waiting times**

- Ideally, we should also accord priority to the more severe cases
- ...But let's do one step at a time

# Waiting Time

We can **build waiting time variables**, similar to what we did for the makespan

```
def add_waittime_variables mdl, levels, codes_by_idx, tasks, eoh):
    codes = ['red', 'yellow', 'green', 'white']
    fbc = {c: [] for c in codes}
    for idx in levels.keys():
        fbc[codes_by_idx[idx]].append(tasks[idx, 0, 0])

    obj_by_code = {}
    for code in codes:
        obj_by_code[code] = mdl.NewIntVar(0, len(codes_by_idx)*eoh, f'wt_{code}')
        if len(fbc[code]) > 0:
            mdl.Add(obj_by_code[code] == sum(t.start for t in fbc[code]))
    obj_by_code['all'] = mdl.NewIntVar(0, len(codes_by_idx) * eoh, 'wt_global')
    mdl.Add(obj_by_code['all'] == sum(obj_by_code[c] for c in codes))
    return obj_by_code
```

# Waiting Time

First, we **collect the first task** for each patient, by code:

```
def add_waittime_variables mdl, levels, codes_by_idx, tasks, eoh):  
    ...  
    fbc = {c: [] for c in codes}  
    for idx in levels.keys():  
        fbc[codes_by_idx[idx]].append(tasks[idx, 0, 0])  
    ...
```

Then, we **build and constrain** the waiting time variables

```
def add_waittime_variables mdl, levels, codes_by_idx, tasks, eoh):  
    ...  
    obj_by_code = {}  
    for code in codes:  
        obj_by_code[code] = mdl.NewIntVar(0, len(codes_by_idx)*eoh, f'wt_{code}')  
        if len(fbc[code]) > 0:  
            mdl.Add(obj_by_code[code] == sum(t.start for t in fbc[code]))  
    obj_by_code['all'] = mdl.NewIntVar(0, len(codes_by_idx) * eoh, f'wt_globai')
```

# Waiting Time

We can formulate an **updated problem model**:

```
In [2]: npatients = 5
mdl = cp_model.CpModel()
levels, codes_by_idx = er.build_levels(fdata.iloc[:npatients])
eoh = er.get_horizon(levels)
rl, dl = {idx:0 for idx in levels}, {idx:eoh for idx in levels}
tasks, last = er.add_all_vars(mdl, levels, rl, dl)
er.add_all_precedences(mdl, levels, tasks)
er.add_cumulatives(mdl, tasks, capacities)
er.add_all_no_overlap(mdl, levels, tasks)
obj_by_code = er.add_waittime_variables(mdl, levels, codes_by_idx, tasks, eoh)
mdl.Minimize(obj_by_code['all'])
slv = cp_model.CpSolver()
slv.parameters.max_time_in_seconds = 10
status = slv.Solve(mdl)
er.print_outcome(slv, levels, tasks, codes_by_idx, status)
```

Solver status: optimal, time(CPU sec): 0.01, objective: 2.0

0(green): visit(0-1), RX(1-3), visit(3-4)  
1(green): visit(0-1), lab(1-5), visit(5-6)  
2(white): visit(0-1), otolaryngological visit(1-5), visit(5-6)  
6(green): visit(1-2), ultrasound(2-4), RX(4-6), lab(6-10), CT scan(10-14), visit(14-15), CT scan(15-19), visit(19-20)  
7(green): visit(1-2)


# Evaluation

We will now try to **compare the complexity** with that of the previous one

```
In [5]: def solve_basic_waittime_problem(npatients, capacities):
        mdl = cp_model.CpModel()
        levels, codes_by_idx = er.build_levels(fdata.iloc[:npatients])
        eoh = er.get_horizon(levels)
        rl, dl = {idx:0 for idx in levels}, {idx:eoh for idx in levels}
        tasks, last = er.add_all_vars(mdl, levels, rl, dl)
        er.add_all_precedences(mdl, levels, tasks)
        er.add_cumulatives(mdl, tasks, capacities)
        er.add_all_no_overlap(mdl, levels, tasks)
        obj_by_code = er.add_waittime_variables(mdl, levels, codes_by_idx, tasks, eoh)
        mdl.Minimize(obj_by_code['all'])
        slv = cp_model.CpSolver()
        slv.parameters.max_time_in_seconds = 10 # time limit
        status = slv.Solve(mdl)
        return status, slv.UserTime()
```


# Evaluation

```
In [12]: tlist, plist = [], range(10, 20, 3)
for npt in plist:
    s, t = solve_basic_waittime_problem(npt, capacities)
    tlist.append(t)
er.plot_scalability_evaluation(plist, np.log10(tlist), figsize=figsize)
```



# Evaluation

```
In [12]: tlist, plist = [], range(10, 20, 3)
         for npt in plist:
             s, t = solve_basic_waittime_problem(npt, capacities)
             tlist.append(t)
         er.plot_scalability_evaluation(plist, np.log10(tlist), figsize=figsize)
```



Sum-based objectives are much harder in CP and max-based objectives



# Goal Programming

---

# Goal Programming

**Any practical application of our method faces two big issues**

First, we have limited scalability

- The problem is computationally challenging
- Note that we may still be getting good solutions quite early
- ...But we need to compare multiple solution approaches to know that

Second, we need to handle patient priorities

- We need to make sure that (e.g.) red-coded patients are handled first
- ...Then yellow, green, and finally white patients

**We can actually address both with a single technique!**

(At least to some degree)

# Goal Programming

Goal Programming is a strategy for handling **lexicographic costs**

- Let's consider a problem with **multiple cost functions**  $f_0(x), f_1(x), \dots$
- Let  $f_i$  be **strictly more important** than  $f_j$  iff  $i < j$

Then **Goal Programming** refers to this simple algorithm:

- Optimize  $f_0(x)$  to obtain its optimal value  $z_0^*$
- Post a new constraint in the model in the form  $f_0(x) \leq z_0^*$
- Move to the next cost function and repeat

**The approach is very simple, but it works well in practice**

A typical example: the Capacitated Vehicle Routing Problem

- $f_0(x)$  = number of vehicles,  $f_1(x)$  = length of the routes
- Thanks to GP, we get an upper bound on each cost at each step
- ...Which may provide a computational advantage

# Implementing Goal Programming

As a first step, we need to handle **upper bounds** on the waiting times

```
def solve_bounded_waittime_problem(levels, codes_by_idx, codes, capacities,
    ub_by_code={}, tlim=None):
    mdl = cp_model.CpModel()
    eoh = get_horizon(levels)
    rl, dl = {idx:0 for idx in levels}, {idx:eoh for idx in levels}
    tasks, last = add_all_vars(mdl, levels, rl, dl)
    add_all_precedences(mdl, levels, tasks)
    add_cumulatives(mdl, tasks, {'visit': 3, 'ultrasound': 2, 'RX': 2})
    add_all_no_overlap(mdl, levels, tasks)
    obj_by_code = add_waittime_variables(mdl, levels, codes_by_idx, tasks, eoh)

    for code in ub_by_code: # <-- Enforce upper bounds
        mdl.Add(obj_by_code[code] <= ub_by_code[code])

    mdl.Minimize(obj_by_code[codes[-1]]) # <-- Focus by default on the last code
    slv = cp_model.CpSolver()
    if tlim is not None: slv.parameters.max_time_in_seconds = tlim
```

# Implementing Goal Programming

Then we need to code the approach in a function

```
def goal_programming(levels, codes_by_idx, codes, capacities, tlim=None, verbose=1):
    wt_by_code, ttime = {}, 0
    for i, code in enumerate(codes):
        l_tlim = None if tlim is None else (tlim-ttime) / (len(codes)-i)
        status, slv, tasks = solve_bounded_waittime_problem(levels, codes_by_idx,
                                                             codes[:i+1], capacities, ub_by_code=wt_by_code, tlim=l_tlim)
        ttime += slv.UserTime()
        if verbose:
            print_outcome(slv, levels, tasks, codes_by_idx, status, (i == len(codes)-1))
        if status in (cp_model.OPTIMAL, cp_model.FEASIBLE):
            wt_by_code[code] = int(slv.ObjectiveValue())
            starts = {tidx:slv.Value(tasks[tidx].start) for tidx in tasks}
        else:
            wt_by_code, starts = None, None
            break
    if wt_by_code is not None:
```

# Implementing Goal Programming

## Finally, we can test the approach

We choose a **priority for our codes**, then we call the `goal_programming` function

```
In [8]: codes = ['red', 'yellow', 'green', 'white']
npatients = 10
levels, codes_by_idx = er.build_levels(fdata.iloc[:npatients])
_, _, _ = er.goal_programming(levels, codes_by_idx, codes, capacities, tlim=10)
```

```
Solver status: optimal, time(CPU sec): 0.01, objective: 0.0
```

```
Solver status: optimal, time(CPU sec): 0.01, objective: 0.0
```

```
Solver status: optimal, time(CPU sec): 0.02, objective: 9.0
```

```
Solver status: optimal, time(CPU sec): 0.04, objective: 3.0
```

```
0(green): visit(0-1), RX(1-3), visit(3-4)
```

```
9(green): visit(0-1)
```

```
12(red): visit(0-1), RX(3-5), visit(5-6)
```

```
6(green): visit(1-2), ultrasound(2-4), RX(4-6), lab(6-10), CT scan(10-14), visit(14-15), CT scan(15-19), visit(19-20)
```

```
7(green): visit(1-2)
```

```
8(green): visit(1-2), RX(2-4), prescription(4-8), visit(8-9)
```

```
1(green): visit(2-3), lab(3-7), visit(7-8)
```

```
10(green): visit(2-3), lab(3-7), CT scan(7-11), visit(11-12)
```


```
11(green): visit(2-3), ultrasound(3-5), lab(5-9), prescription(9-13), visit(13-14)
```

```
2(white): visit(3-4), otolaryngological visit(4-8), visit(8-9)
```

# Evaluation

We can now evaluate the scalability of the approach

```
In [13]: tlist, olist, plist = [], [], range(10, 20, 3)
for npt in plist:
    levels, codes_by_idx = er.build_levels(fdata.iloc[:npt])
    t, o, _ = er.goal_programming(levels, codes_by_idx, codes, capacities, tlim=10, verbose=0)
    tlist.append(t)
    olist.append(o['all'])
er.plot_scalability_evaluation(plist, np.log10(tlist), figsize=figsize)
```



# Warm Starting

A drawback of GP is that search restarts from scratch at each iteration

- This can be mitigated by **warm starting** the solver
- In Google ortools, this is done via so-called "hints"

```
def goal_programming(levels, codes_by_idx, codes, capacities,
    tlim=None, verbose=1, hints=None):
    ...
    status, slv, tasks = solve_bounded_waittime_problem(levels, codes_by_idx,
        codes[:i+1], capacities, ub_by_code=wt_by_code,
        tlim=l_tlim, hints=hints) # <-- here we pass hints
    ...
    if status in (cp_model.OPTIMAL, cp_model.FEASIBLE):
        ...
        starts = {tidx:slv.Value(tasks[tidx].start) for tidx in tasks}
        hints = starts # <-- hints are the start in the last solution
    ...
```



# Warm Starting

And finally we need to take into account the hints in the model:

```
def solve_bounded_waittime_problem(levels, codes_by_idx, codes, capacities,
    ub_by_code={}, tlim=None, hints=None):
    ...
    # Add hints
    if hints is not None:
        for (idx, k, i), stval in hints.items():
            mdl.AddHint(tasks[idx,k,i].start, stval)
    ...
```


- The solver will **try to assign** to each value the corresponding hint
- In case this proves impossible, it will search as usual
- If properly done, this guarantees that we start from a known feasible solution

This is usually referred to as **warm starting** in the optimization field

# Warm Starting

## Let's test this updated implementation

```
In [14]: tlist2, plist2 = [], range(10, 23, 3)
for npt in plist2:
    levels, codes_by_idx = er.build_levels(fdata.iloc[:npt])
    t, _, _ = er.goal_programming(levels, codes_by_idx, codes, capacities, tlim=10, verbose=0)
    tlist2.append(t)
er.plot_scalability_evaluation(plist2, np.log10(tlist2), figsize=figsize)
```



# Large Neighborhood Search

---

# Large Neighborhood Search

## Exact optimization approaches (e.g. classical CP/SMT)...

- ...Enable **convenient modeling** and are very effective on small problems
- ...But have trouble dealing with large-scale optimization problems

## Heuristic approaches (e.g. local search)...


- ...Can deal with large scale problems
- ...But have trouble escaping local minima

## The two approaches can be combined to obtain **Large Neighborhood Search**

- We operate by trying to improve an **incumbent solution**
- Improving solutions are searched in a **large neighborhood**
- Each neighborhood is explored via a complete search on a **restricted problem**
- ...Typically: we fix all variables, except a small subset (**relaxed variables**)

# Large Neighborhood Search

Here's a visual scheme of how the approach works



**LNS combines the strengths of its component approaches**

- It can efficiently handle large-scale problems (like local search)
- It can escape local minima thanks to the use of large neighborhoods
- It retains the modeling flexibility of CP, SMT, MILP, etc.

**Choosing which variables to relax at each iteration...**


- ...Can be a major design decision, and can be done using domain knowledge
- ...But actually even a simple **random selection** often works very well

# Variable Fixing in Scheduling

LNS can be very effective on **scheduling problems**...

- ...But it is not straightforward to implement
- ...Since the classical "variable" fixing approach provides **too little flexibility**

**Consider the following schedule for a resource with capacity 2:**



- Say the we choose to relax the green tasks, and keep the yellow ones fixed
- ...Then we still cannot make any improvement!


**We need to retain some ability to **shift tasks in time****

- But we still need to keep them somewhat "fixed"
- ...Or there is no point in defining a neighborhood

# Partial Order Schedule

The solution is converting a fixed-start schedule into a **Partial Order Schedule**

A Partial Order Schedule is an **augmented task graph**




- The graph contains all the original precedence relations (in this case none)
- ...Plus additional precedences, introduced to **prevent resource conflicts**

## In the figure:

- As long as the black arcs (precedences) are respected
- ...Tasks can be moved **without any resource conflict**

# Partial Order Scheduling

A POS can be obtained by solving a **minimum flow problem**




- First, we introduce a fake "source" and "sink" task
- Then, we build an arc for each pair of tasks such that  $e_i \leq s_j$ 
  - The source can be connected only to the first tasks
  - ...And the sink has an incoming arc for each last task
- Then we say each task "requires" flow equal to its resource demand
  - In this example, all requirements are 1



# Partial Order Scheduling

A POS can be obtained by solving a minimum flow problem



- Finally, we route the minimum amount of flow that satisfies the requirements
  - ...For example using some variant of the Ford-Fulkerson's method
  - We will not look into details into this
- The arcs with non-zero flow become part of the POS
  - In the figure, these are the arcs in black
  - You can check that the flow is actually minimum

# POS Conversion - Implementation

Since we have only unary demands, our implementation can be simpler

First, let's prepare a function to collect and sort the tasks for each resource:

```
def sol_to_pos(levels, starts, capacities):
    aplus, aminus = {}, {}
    tasks = {r:[] for r in capacities}
    for idx, k, i in starts:
        ttype = levels[idx][k][i]
        if ttype in capacities:
            tasks[ttype].append((idx, k, i))
            aplus[idx, k, i], aminus[idx, k, i] = [], []
    for res in capacities:
        tasks[res] = sorted(tasks[res], key=lambda t: starts[t]) # <-- sort by start
    for res, cap in capacities.items():
        for _ in range(cap):
            tasks[res] = unit_flow(levels, tasks[res], starts, aplus, aminus)
    return aplus, aminus
```

# POS Conversion - Implementation

We store the POS as a collection of outgoing and ingoing arcs

This is called a **forward and backward star** representation:

```
def sol_to_pos(levels, starts, capacities):
    aplus, aminus = {}, {}
    ...
    for idx, k, i in starts:
        ...
        if ttype in capacities:
            ...
            aplus[idx, k, i], aminus[idx, k, i] = [], []
    ...
    return aplus, aminus
```

- The chosen representation allows for efficient graph traversal
- It also makes it easy to detect arcs implied via the transitive property

# POS Conversion - Implementation

Then, we collect the tasks for each resource:

```
def sol_to_pos(levels, starts, capacities):  
    ...  
    tasks = {r:[] for r in capacities}  
    for idx, k, i in starts:  
        ttype = levels[idx][k][i]  
        if ttype in capacities:  
            tasks[ttype].append((idx, k, i))  
        ...  
    for res in capacities:  
        tasks[res] = sorted(tasks[res], key=lambda t: starts[t])  
    ...
```

- We sort each collection by increasing start time

# POS Conversion - Implementation

Finally, we route single units of flows

```
def sol_to_pos(levels, starts, capacities):  
    ...  
    for res, cap in capacities.items():  
        for _ in range(cap):  
            stasks[res] = unit_flow(levels, stasks[res], starts, aplus, aminus)  
    return aplus, aminus
```

- Each unit corresponds to one unit of resource capacity
- The flow represents how the resource is "passed" from one task to the next
- We do not build explicitly the possible arcs: they are implied by the start times

# POS Conversion - Implementation

Here is the algorithm for routing each flow unit

```
def unit_flow(levels, stasks, starts, aplus, aminus):
    nonprocessed = []
    src, dur = None, 0
    for dst in stasks:
        if src is not None and starts[src] + dur <= starts[dst]:
            aplus[src].append(dst) # Build an arc (forward star)
            aminus[dst].append(src) # Build an arc (backward star)
            src = None # Reset the source
        if src is None:
            src = dst # Set a new source
            dur = get_dur(levels[src[0]][src[1]][src[2]])
        else:
            nonprocessed.append(dst) # Store as non-processed
    return nonprocessed
```

# POS Conversion - Implementation

## Let's test the function

```
In [15]: codes = ['red', 'yellow', 'green', 'white']
npatients = 3
levels, codes_by_idx = er.build_levels(fdata.iloc[:npatients])
_, _, starts = er.goal_programming(levels, codes_by_idx, codes, capacities, tlim=10)
aplus, aminus = er.sol_to_pos(levels, starts, capacities)
aplus
```

```
Solver status: optimal, time(CPU sec): 0.00, objective: 0.0
Solver status: optimal, time(CPU sec): 0.00, objective: 0.0
Solver status: optimal, time(CPU sec): 0.00, objective: 0.0
Solver status: optimal, time(CPU sec): 0.01, objective: 0.0
```

```
0(green): visit(0-1), RX(1-3), visit(3-4)
1(green): visit(0-1), lab(1-5), visit(5-6)
2(white): visit(0-1), otolaryngological visit(1-5), visit(5-6)
```

```
Out[15]: {(0, 0, 0): [(0, 2, 0)],
(0, 1, 0): [],
(0, 2, 0): [(1, 2, 0)],
(1, 0, 0): [(2, 2, 0)],
(1, 2, 0): [],
(2, 0, 0): [],
(2, 2, 0): []}
```

# POS Relaxation

## A POS corresponds to a set of feasible schedules

- Tasks can be freely shifted, and no resource constraints is be violated
- There are **no longer ordering decisions** to be made
- Typically, choosing one of the possible schedules becomes a poly-time problem

However, this means that we cannot really make significant changes

## For using a POS in LNS, we need to **relax** it

Typically:

- We choose a set of tasks to relax
- ...And we disconnect them from the precedence network
  - First, we connect all task predecessors to the task successors
  - Then, we remove all the additional POS arcs linked to the relaxed tasks



# POS Relaxation - Implementation

Here's an implemented algorithm for removing/relaxing POS tasks

```
def remove_task_from_pos(aplus, aminus, task_key):  
    # Transfer transitive arcs  
    for src in aminus[task_key]:  
        for dst in aplus[task_key]:  
            aminus[dst].append(src)  
            aplus[src].append(dst)  
    # Remove ingoing arcs  
    for src in aminus[task_key]:  
        aplus[src].remove(task_key)  
    # Remove outgoing arcs  
    for dst in aplus[task_key]:  
        aminus[dst].remove(task_key)  
    # Remove the node from the arc  
    del aplus[task_key]  
    del aminus[task_key]
```

# POS Relaxation - Implementation

In our LNS approach we will relax all tasks related to selected patients

```
def remove_patient_from_pos(aplus, aminus, target_idx):  
    aplus_res = copy.deepcopy(aplus)  
    aminus_res = copy.deepcopy(aminus)  
    for idx, k, i in aplus:  
        if idx == target_idx:  
            remove_task_from_pos(aplus_res, aminus_res, (idx, k, i))  
    return aplus_res, aminus_res
```

```
In [16]: aplus2, aminus2 = er.remove_patient_from_pos(aplus, aminus, 2)  
         aplus2
```

```
Out[16]: {(0, 0, 0): [(0, 2, 0)],  
          (0, 1, 0): [],  
          (0, 2, 0): [(1, 2, 0)],  
          (1, 0, 0): [],  
          (1, 2, 0): []}
```

# Restricted Problem

We now need to take into account the additional precedences in the model

```
def solve_bounded_waittime_problem(levels, codes_by_idx, codes, capacities,
    ub_by_code={}, tlim=None, hints=None, aplus=None):
    ...
    # Add extra precedences
    if aplus is not None:
        for (idx, k, i), out_arcs in aplus.items():
            for idx2, k2, i2 in out_arcs:
                mdl.Add(tasks[idx,k,i].end <= tasks[idx2,k2,i2].start)
    ...
```

Even if we still need to assign **all the start times**...

- Precedence constraints are efficient to propagate
- They are effective at narrowing the variable domains
- ...And in this case they prevent many resource conflicts

**If many POS constraints are added, the problem is much easier to solve**

# LNS - Initial Solution

## We are ready to define our LNS approach

First, we need an initial solution:

```
In [17]: init_time = 5 # Time limit for the initial solution
         codes = ['red', 'yellow', 'green', 'white']
         npatients = 100

         levels, codes_by_idx = er.build_levels(fdata.iloc[:npatients])

         ttime, wt_by_code, starts = er.goal_programming(levels, codes_by_idx, codes, capacities=capacities,
                                                         tlim=init_time, verbose=0)
         print(f'Initial solution in {ttime:.2f} sec, {wt_by_code}')
```

Initial solution in 5.01 sec, {'red': 5, 'yellow': 577, 'green': 1209, 'white': 10, 'all': 1801}

- We can obtain this one as usual
- It's better if it's of decent quality
- ...But we don't want to spend more than a few seconds on this

# LNS - Baseline

As a baseline, let's see the solution quality we can obtain in 30 seconds

```
In [22]: init_time = 30 # Time limit for the initial solution
        codes = ['red', 'yellow', 'green', 'white']
        npatients = 100

        levels, codes_by_idx = er.build_levels(fdata.iloc[:npatients])

        ttime, wt_by_code, starts = er.goal_programming(levels, codes_by_idx, codes, capacities=capacities,
                                                         tlim=init_time, verbose=0)
        print(f'Initial solution in {ttime:.2f} sec, {wt_by_code}')
```

Initial solution in 30.02 sec, {'red': 5, 'yellow': 307, 'green': 1343, 'white': 22, 'all': 1677}

# LNS - Upper bounds

In LNS we will repeatedly be seeking for an **improving** solution

...So we need to modify our G.P. loop to take **upper bounds** into account:

```
def goal_programming(levels, codes_by_idx, codes, capacities,
    tlim=None, verbose=1, hints=None, ub_by_code=None):
    # Handle the upper bounds
    if ub_by_code is None:
        wt_by_code = {}
    else:
        wt_by_code = copy.deepcopy(ub_by_code)
    ...
    return ttime, wt_by_code, starts
```

- The bounds are then passed to the function that builds and solves the model
- ...Which was already designed to deal with the bounds from Goal Programming

# LNS - Main Loop

Here is the main structure of our LNS implementation:

```
def scheduling_lns(levels, codes_by_idx, codes, capacities,
                  tlim, init_time, it_time, nb_size, verbose=1):
    ... # Build initial solution
    while tlim - ttime > 0:
        aplus, aminus = sol_to_pos(levels, starts, capacities) # Obtain a POS
        relaxed = np.random.choice(patients, nb_size, replace=False) # Relax patients
        for idx in relaxed: aplus, aminus = remove_patient_from_pos(aplus, aminus, idx)
        ub_by_code = copy.deepcopy(wt_by_code) # Require an improvement
        ub_by_code['all'] -= 1
        ... # Re-solve
        ttime += max(itime, 0.1) # Update the time limit
        ... # Update the best solution
```

The full implementation can be found in our support module

# Evaluation

**Time to test! Waiting times are sorted by code, the last is the total**

```
In [23]: # Configuration
tlim = 30
init_time, it_time, nb_size = 3, 1, 4
codes = ['red', 'yellow', 'green', 'white']
npatients = 100
levels, codes_by_idx = er.build_levels(fdata.iloc[:npatients])

er.scheduling_lns(levels, codes_by_idx, codes, capacities, tlim, init_time, it_time, nb_size);
```

Initial solution in 2.90 sec, {'red': 5, 'yellow': 649, 'green': 1386, 'white': 10, 'all': 2050}

Total time: 3.67, neighborhood explored in 0.77 sec, waiting times 5,630,1364,10,2009

Total time: 4.45, neighborhood explored in 0.78 sec, waiting times 5,625,1364,10,2004

Total time: 5.20, neighborhood explored in 0.75 sec, waiting times 5,622,1364,10,2001

Total time: 5.98, neighborhood explored in 0.78 sec, waiting times 5,613,1361,10,1989

Total time: 6.63, neighborhood explored in 0.65 sec, waiting times 5,605,1359,10,1979

Total time: 7.28, neighborhood explored in 0.66 sec, waiting times 5,590,1338,10,1943

Total time: 8.04, neighborhood explored in 0.76 sec, waiting times 5,583,1333,10,1931

Total time: 8.82, neighborhood explored in 0.77 sec, waiting times 5,567,1330,10,1912

Total time: 9.60, neighborhood explored in 0.78 sec, waiting times 5,557,1330,10,1902

Total time: 10.25, neighborhood explored in 0.65 sec, waiting times 5,549,1328,10,1892

Total time: 11.03, neighborhood explored in 0.78 sec, waiting times 5,529,1324,10,1868

Total time: 11.66, neighborhood explored in 0.62 sec, waiting times 5,508,1322,10,1845

Total time: 12.14, neighborhood explored in 0.48 sec, waiting times 5,497,1315,10,1827

Total time: 12.92, neighborhood explored in 0.78 sec, waiting times 5,487,1305,10,1807

Total time: 13.68, neighborhood explored in 0.76 sec, waiting times 5,482,1305,10,1802



# Considerations

## **LNS is the industrial optimizer's secret weapon**

- It's flexible, robust, scalable, and not too difficult to implement
  - Except for scheduling problems, but you've now seen how to do it ;-)
- It's usually possible to do better than LNS, but it's also considerably harder

## **Goal programming is not only about lexicographic costs**

- It's about simplifying a problem via constraints
- E.g. constraints may actually help in CP
- Sometimes special optimization are possible:
  - Include only some of the variables (e.g. red-codes only)
  - Simplified problem formulation (e.g. fixed number of vehicles)

## **Do not underestimate warm starts**

- Without our little trick at the beginning
- ...The search for an initial solution with 100 patients would fail!